

Lab 5 – Manage Kubernetes Containers

Introduction

In this lab, you will understand and work with **core concepts of kubernetes**, Includes PODs, Labels, Replica Sets, Replication Controllers, Deployments, Services, Daemons and Namespaces.

1. Namespaces

1.1 Login to pod0-master.origin.com host as **root** user, if not already:

Copy

```
ssh root@pod0-master.origin.com
```

Kubernetes supports multiple virtual clusters backed by the same physical cluster. These virtual clusters are called namespaces. Within the same namespace, kubernetes objects name should be unique. Different objects in different namespaces may have the same name.

Kubernetes comes with three initial namespaces

- **default:** the default namespace for objects with no other namespace
- **kube-system** the namespace for objects created by the kubernetes system
- **kube-public** The namespace is created automatically and readable by all users.

1.2 To get namespaces, Verify the default existing namespaces

Copy

```
kubectl get namespaces
```

Output:

NAME	STATUS	AGE
default	Active	9h

kube-public	Active	9h
kube-system	Active	9h

1.3 Create a new namespaces as “kube-core”

Copy

```
kubectl create namespace kube-core
```

Output:

```
namespace "kube-core" created
```

1.4 List all namespaces

Copy

```
kubectl get namespaces
```

Output:

NAME	STATUS	AGE
default	Active	9h
kube-core	Active	9h
kube-public	Active	9h
kube-system	Active	9h

1.5 Describe the specified namespace:

Copy

```
kubectl describe namespace kube-core
```

Output:

Name: kube-core

Labels: <none>

Annotations: <none>

Status: Active

No resource quota.

No resource limits.

1.6 Verify the labels for namespaces

Copy

```
kubectl get namespaces --show-labels
```

Output:

NAME	STATUS	AGE	LABELS
default	Active	34m	<none>
kube-core	Active	21s	<none>
kube-public	Active	34m	<none>
kube-system	Active	34m	<none>

1.7 Verify the contexts for the current namespaces

Copy

```
kubectl config view
```

Output:

```
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: REDACTED
    server: https://10.1.64.178:6443
  name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubernetes-admin
  name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data: REDACTED
    client-key-data: REDACTED
```

1.8 Set the context for **kube-core** namespaces

Copy

```
kubectl config set-context production --namespace=kube-core --cluster=kubernetes --user=admin-user
```

Output:

```
Context "production" created.
```

1.9 List the available contexts

Copy

```
kubectl config get-contexts
```

Output:

CURRENT NAMESPACE	NAME	CLUSTER	AUTHINFO
*	kubernetes-admin@kubernetes	kubernetes	kubernetes-admin
kube-core	production	kubernetes	admin-user

1.10 Switch to the newly created context

Copy

```
kubectl config use-context production
```

Output:

```
Switched to context "production".
```

1.11 Verify the current context in use

Copy

```
kubectl config current-context
```

Output:

```
production
```

1.12 Display merged kubeconfig settings

Copy

```
kubectl config view
```

Output:

```
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: REDACTED
    server: https://10.1.64.178:6443
  name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubernetes-admin
  name: kubernetes-admin@kubernetes
- context:
    cluster: kubernetes
```

```
namespace: kube-core

user: admin-user

name: production

current-context: production

kind: Config

preferences: {}

users:

- name: kubernetes-admin

  user:

    client-certificate-data: REDACTED

    client-key-data: REDACTED
```

1.13 Switch back to the default namespaces:

Copy

```
kubectl config use-context kubernetes-admin@kubernetes
```

Output:

```
Switched to context "kubernetes-admin@kubernetes".
```

2. User in Kubernetes

In Kubernetes there are **two categories** of users:

- Service accounts managed by Kubernetes
- Normal users – Outside of kubernetes

An admin distributing **private keys**, a user store like Keystone or Google Accounts, even a file with a list of usernames and passwords.

Service accounts are users managed by the Kubernetes API. They are bound to **specific namespaces**, and created automatically by the API server or manually through API calls. Service accounts are tied to a set of credentials stored as Secrets, which are mounted into pods allowing in-cluster processes to talk to the Kubernetes API.

a. Create Service Account

2.1 We are creating a Service Account with the name admin-user in namespace kube-core.

Copy

```
cat > credadmin.yaml <<EOF

apiVersion: v1

kind: ServiceAccount

metadata:

  name: admin-user

  namespace: kube-core

EOF
```

2.2 Create the service account for **kube-core** namespace

Copy

```
kubectl create -f credadmin.yaml
```

Output:

```
serviceaccount "admin-user" created
```

b. Create ClusterRoleBinding

2.3 In most cases after provisioning our cluster using **kops** or **kubeadm** or any other popular tool admin Role already exists in the cluster. We can use it and create only **RoleBinding** for our **ServiceAccount**.

Copy

```
cat > clusterrolebinding.yaml <<EOF

apiVersion: rbac.authorization.k8s.io/v1beta1

kind: ClusterRoleBinding

metadata:

  name: admin-user

roleRef:

  apiGroup: rbac.authorization.k8s.io

  kind: ClusterRole

  name: cluster-admin

subjects:

- kind: ServiceAccount

  name: admin-user

  namespace: kube-core

EOF
```

2.4 Create a **ClusterRoleBinding** for service account

Copy

```
kubectl create -f clusterrolebinding.yaml
```

Output:

```
clusterrolebinding.rbac.authorization.k8s.io "admin-user" created
```

c. Cleanup Namespaces

2.5 Delete the context

Copy

```
kubectl config delete-context production
```

Output:

```
deleted context production from /root/.kube/config
```

2.6 Delete the namespaces

Copy

```
kubectl delete namespaces kube-core
```

Output:

```
namespace "kube-core" deleted
```

3. PODs

In Kubernetes, a group of one or more containers is called a **pod**. Containers in a pod are deployed together, and are **started**, **stopped**, and **replicated** as a group.

3.1 Create a nginx pod, name as “**pod-nginx.yaml**” file.

Copy

```
cat > pod-nginx.yaml <<EOF
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:

  name: mynginx

  namespace: default

  labels:

    run: nginx

spec:

  containers:

  - name: mynginx

    image: nginx:latest

    ports:

      - containerPort: 80

EOF
```

A pod definition is a declaration of a **desired state**.

- **Desired state** is a very important concept in the Kubernetes model. Many things present a desired state to the system, and it is Kubernetes' responsibility to make sure that the current state matches the desired state.

3.2 Create a pod containing an nginx server from the **pod-nginx.yaml** file

Copy

```
kubectl create -f pod-nginx.yaml
```

Output:

```
pod "mynginx" created
```

3.3 List all pods are running state:

Copy

```
kubectl get pods -o wide
```

Output:

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
mynginx e.onecloud.com	1/1	Running	0	2m	20.1.1.3	pod0-nod

Note: Wait till the pod status changes to “**running**” status. **Re-run** the above command to verify of it.

A pod can be in one of the following phases:

- **Pending:** the API Server has created a pod resource and stored it in etcd, but the pod has not been scheduled yet, nor have container images been pulled from the registry.
- **Running:** the pod has been scheduled to a node and all containers have been created by the kubelet.
- **Succeeded:** all containers in the pod have terminated successfully and will not be restarted.
- **Failed:** all containers in the pod have terminated and, at least one container has terminated in failure.
- **Unknown:** The API Server was unable to query the state of the pod, typically due to an error in communicating with the kubelet.

3.4 Describe the created **nginx** pod in detailed:

Copy

```
kubectl describe pod mynginx
```

Output:

```
Name:          mynginx
Namespace:     default
Node:          pod30-node.onecloud.com/10.1.64.178
```

Start Time: Sun, 22 Apr 2018 08:33:21 +0000

Labels: run=nginx

Annotations: <none>

Status: Running

IP: 20.1.1.3

Containers:

....

....

Events:

Type	Reason	Age	From
Message			
----	-----	----	----

Normal	Scheduled	1m	default-scheduler
Successfully assigned mynginx to pod30-node.onecloud.com			
Normal	SuccessfulMountVolume	1m	kubelet, pod30-node.onecloud.com
MountVolume.SetUp succeeded for volume "default-token-nf7p8"			
Normal	Pulling	1m	kubelet, pod30-node.onecloud.com
pulling image "nginx:latest"			
Normal	Pulled	48s	kubelet, pod30-node.onecloud.com
Successfully pulled image "nginx:latest"			
Normal	Created	48s	kubelet, pod30-node.onecloud.com
Created container			
Normal	Started	48s	kubelet, pod30-node.onecloud.com
Started container			

4. Labels

In Kubernetes, labels are a system to organize **objects** into **groups**. Labels are **key-value pairs** that are attached to each object.

Label selectors can be passed along with a request to the apiserver to retrieve a list of objects which match that label selector.

4.1 Create a label to a pod, add a labels section under metadata in the pod definition:

Copy

```
cat pod-nginx.yaml
```

4.2 Label the **mynginx** pod

Copy

```
kubectl label pod mynginx type=webserver
```

Output:

```
pod "mynginx" labeled
```

4.3 List the pods based on specified label

Copy

```
kubectl get pods -l type=webserver
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
mynginx	1/1	Running	0	4m

4.4 Labels can be applied not only to pods but also to other Kubernetes objects like nodes.

Copy

```
kubectl get nodes
```

Output:

NAME	STATUS	ROLES	AGE	VERSION
pod0-master.onecloud.com	Ready	master	8h	v1.10.0
pod0-node.onecloud.com	Ready	node	8h	v1.10.0

4.5 Apply the label to **pod0-master.onecloud.com** node

Copy

```
kubectl label node pod0-master.onecloud.com rack=rack01
```

Output:

```
node "pod0-master.onecloud.com" labeled
```

4.6 Verify that the **master** node is labeled

Copy

```
kubectl get nodes -l rack=rack01
```

Output:

NAME	STATUS	ROLES	AGE	VERSION
pod0-master.onecloud.com	Ready	master	8h	v1.10.0

4.7 Apply label to **pod0-master.onecloud.com** node

Copy

```
kubectl label node pod0-master.onecloud.com rack=rack01
```

Output:

```
node "pod0-master.onecloud.com" labeled
```

4.8 Verify that the worker node is labeled

Copy

```
kubectl get nodes -l rack=rack01
```

Output:

NAME	STATUS	ROLES	AGE	VERSION
pod38-node.onecloud.com	Ready	master	8h	v1.10.0
pod38-node.onecloud.com	Ready	node	8h	v1.10.0

Labels are also used as selector for services and deployments.

5. Replica Controllers

A Replica Controller ensures that a specified number of pod replicas are running at any one time. In other words, a Replica Controller makes sure that a pod or homogeneous set of pods are always up and available. If there are too many pods, it will kill some. If there are too few, it will start more. Unlike manually created pods, the pods maintained by a Replica Controller are automatically replaced if they fail, get deleted, or are terminated.

A Replica Controller configuration consists of:

- The number of replicas desired
- The pod definition
- The selector to bind the managed pod

A selector is a label assigned to the pods that are managed by the replica controller. Labels are included in the pod definition that the replica controller instantiates. The replica controller uses the selector to determine how many instances of the pod are already running in order to adjust as needed.

5.1 Create a replica controller with **replica** of 3 for our **nginx** pod, and name as “**nginx-rc.yaml**” file:

Copy

```
cat > nginx-rc.yaml <<EOF
---
apiVersion: v1
kind: ReplicationController
metadata:
  name: nginx-rc
  namespace: default
spec:
  replicas: 3
  selector:
    run: nginx
  template:
    metadata:
      name: nginx-rc
      labels:
        run: nginx
    spec:
      containers:
        - name: nginx-rc
          image: nginx:latest
```

```
    ports:

    - containerPort: 80

EOF
```

5.2 Create a replica controller using **nginx-rc.yaml** file:

Copy

```
kubectl create -f nginx-rc.yaml
```

Output:

```
replicationcontroller "nginx-rc" created
```

5.3 List and describe a replica controller

Copy

```
kubectl get rc
```

Output:

NAME	DESIRED	CURRENT	READY	AGE
nginx-rc	3	3	3	1m

Note: Wait till READY count changes to “3”.

5.4 Describe the replica controller which you just created:

Copy

```
kubectl describe rc nginx-rc
```

Output:

```
Name:          nginx-rc
```

```
Namespace:    default
Selector:     run=nginx
Labels:       run=nginx
Annotations:  <none>
Replicas:     3 current / 3 desired
Pods Status:  3 Running / 0 Waiting / 0 Succeeded / 0 Failed
```

...

...

Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	SuccessfulCreate	1m	replication-controller	Created pod: nginx-rc-zpgwx
Normal	SuccessfulCreate	1m	replication-controller	Created pod: nginx-rc-zc6g4

5.5 Verify the pods are in running state:

Copy

```
watch kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
mynginx	1/1	Running	0	7m
nginx-rc-zc6g4	1/1	Running	0	2m

nginx-rc-zpgwx	1/1	Running	0	2m
----------------	-----	---------	---	----

Note: Wait for couple of minutes to get nginx-rc-*** pods to create and READY as **Running** and then press **<ctrl+c>** to interrupt.

5.6 The Replica Controller makes it easy to scale the number of replicas up or down, either manually or by an auto-scaling control agent, by simply updating the replicas field.

Copy

```
kubectl scale rc nginx-rc --replicas=0
```

Output:

```
replicationcontroller "nginx-rc" scaled
```

5.7 Verify that replica controller is scaled **down to 0**

Copy

```
kubectl get rc nginx-rc
```

Output:

NAME	DESIRED	CURRENT	READY	AGE
nginx-rc	0	0	0	7m

5.8 Scale up the replica controller to create new pods

Copy

```
kubectl scale rc nginx-rc --replicas=4
```

Output:

```
replicationcontroller "nginx-rc" scaled
```

5.9 Verify that **nginx-rc** pods **scaled** to 4.

Copy

```
kubectl get rc nginx-rc
```

Output:

NAME	DESIRED	CURRENT	READY	AGE
nginx-rc	4	4	4	9m

Note: Wait till READY count changes to “4”.

Also in case of failure of a node, the replica controller takes care of keep the same number of pods by scheduling the containers running on the failed node to the remaining nodes in the cluster.

5.10 Let us delete one of the pod from the list of running, and you notice the new pod will be created immediately.

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx-rc-fxq2j	1/1	Running	0	2m
nginx-rc-hfgvr	1/1	Running	0	2m
nginx-rc-rdkrq	1/1	Running	0	2m
nginx-rc-rn8rf	1/1	Running	0	2m

Note: Wait for couple of minutes to get pods to create and **Running** in READY state.

Copy

```
kubectl delete pod `kubectl get pods | grep nginx-rc | awk '{print $1}' | head -1`
```

Output:

```
pod "nginx-rc-fxq2j" deleted
```

5.11 Verify that “**nginx-rc-******“, pod is deleted and getting recreated

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx-rc-786x7	0/1	ContainerCreating	0	17s
nginx-rc-hfgvr	1/1	Running	0	3m
nginx-rc-rdkrq	1/1	Running	0	3m
nginx-rc-rn8rf	1/1	Running	0	3m

5.12 Run the below command to delete a replica controller, which deletes all of the pods created.

Copy

```
kubectl delete rc/nginx-rc
```

Output:

```
replicationcontroller "nginx-rc" deleted
```

5.13 Verify the pods are deleted

Copy

```
kubectl get pods
```

Output:

```
No resources found.
```

5.14 Deleting a replica controller deletes all pods managed by that replica. But, because pods created by a replication controller are not actually an integral part of the replication controller, but only managed by it, we can delete only the replication controller and leave the pods running.

Copy

```
kubectl create -f nginx-rc.yaml
```

Output:

```
replicationcontroller "nginx-rc" created
```

5.15 Verify the pod status

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx-rc-hgqv9	1/1	Running	0	1m
nginx-rc-jqh8f	1/1	Running	0	1m
nginx-rc-tk7nc	1/1	Running	0	1m

Note: Wait till the pod ready changes to “**running**”. **Re-run** the above command to verify of it.

5.16 Delete the replication controller with the option **cascade**, which retains the pods by just deleting replication controller.

Copy

```
kubectl delete rc/nginx-rc --cascade=false
```

Output:

```
replicationcontroller "nginx-rc" deleted
```

5.17 Verify the rc pod status

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx-rc-hgqv9	1/1	Running	0	1m
nginx-rc-jqh8f	1/1	Running	0	1m
nginx-rc-tk7nc	1/1	Running	0	1m

5.18 Verify the replication controllers are deleted.

Copy

```
kubectl get rc
```

Output:

```
No resources found.
```

Now there is nothing managing pods, but we can always create a new replication controller with the proper label selector and make them managed again.

5.19 Delete all the pods

Copy

```
kubectl delete pod --all
```

Output:

```
pod "nginx-rc-hgqv9" deleted
```

```
pod "nginx-rc-jqh8f" deleted
```

```
pod "nginx-rc-tk7nc" deleted
```

6. Replica Sets

A Replica Set object is very similar to the replica controller object we practiced before.

6.1 Create a file “**nginx-rs.yaml**” configuration file defines a replica set for the nginx pod

Copy

```
cat > nginx-rs.yaml <<EOF

---

apiVersion: extensions/v1beta1

kind: ReplicaSet

metadata:

  labels:

    run: nginx

  namespace:

  name: nginx-rs

spec:
```

```
replicas: 3

selector:

  matchLabels:

    run: nginx

template:

  metadata:

    labels:

      run: nginx

  spec:

    containers:

      - image: nginx:1.12

        imagePullPolicy: Always

        name: nginx

        ports:

          - containerPort: 80

            protocol: TCP

        restartPolicy: Always

EOF
```

6.2 Create a replica set using **nginx-rs.yaml** file

Copy

```
kubectl create -f nginx-rs.yaml
```

Output:

```
replicaset.extensions "nginx-rs" created
```

6.3 Verify the replicaset status

Copy

```
kubectl get rs -o wide
```

Output:

NAME SELECTOR	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES
nginx-rs 1.12 run=nginx	3	3	3	1m	nginx	nginx:

Note: Wait till READY count changes to “3”.

A replica set behaves exactly like a replication controller, but it has more powerful pod selectors. Whereas a replication controller label selector only allows matching pods that include a certain label, the replica set pod selector also allows matching pods that lack a certain label or pods that include a certain label key, regardless of its value.

6.4 Verify that the status for nginx replicas

Copy

```
kubectl get pods
```

Sample output:

NAME	READY	STATUS	RESTARTS	AGE
nginx-rs-26kw5	1/1	Running	0	6m
nginx-rs-dszmn	1/1	Running	0	6m
nginx-rs-s5cwd	1/1	Running	0	6m

6.5 Delete all the pods

Copy

```
kubectl delete pod --all
```

Sample output:

```
pod "nginx-rs-26kw5" deleted  
pod "nginx-rs-dszmn" deleted  
pod "nginx-rs-s5cwd" deleted
```

6.6 Delete the replicaset

Copy

```
kubectl delete rs nginx-rs
```

Output:

```
replicaset.extensions "nginx-rs" deleted
```

7. Deployments

A Deployment provides declarative updates for pods and replicas. The Deployment object defines the strategy for transitioning between deployments of the same application.

There are basically two ways of updating an application:

- delete all existing pods first and then start the new ones or
- start new ones and once they are up, delete the old ones

The latter, can be done with two different approach:

- add all the new pods and then deleting all the old ones at once
- add new pods at time and then removing old ones one by one

7.1 Create a deployment for our **nginx webserver**, name as **nginx-deploy.yaml** file .

Copy

```
cat > nginx-deploy.yaml <<EOF

apiVersion: extensions/v1beta1

kind: Deployment

metadata:

  generation: 1

  labels:

    run: nginx

  name: nginx

  namespace: default

spec:

  replicas: 2

  selector:

    matchLabels:

      run: nginx

  strategy:

    rollingUpdate:

      maxSurge: 1

      maxUnavailable: 1

    type: RollingUpdate

  template:
```

```
metadata:

  labels:

    run: nginx

spec:

  containers:

  - image: nginx:latest

    imagePullPolicy: Always

    name: nginx

    ports:

      - containerPort: 80

        protocol: TCP

    restartPolicy: Always

EOF
```

7.2 Create a deployment for **nginx-webserver** using **nginx-deploy.yaml** file

Copy

```
kubectl create -f nginx-deploy.yaml
```

Output:

```
deployment.extensions "nginx" created
```

7.3 The deployment creates the following objects

Copy

```
kubectl get all -l run=nginx -o wide
```

Output:

NAME IP	NODE	READY	STATUS	RESTARTS	AGE
nginx-6bfb654d7c-kglvd <none>	pod0-master.onecloud.com	0/1	ContainerCreating	0	16s
nginx-6bfb654d7c-s8k5f <none>	pod0-master.onecloud.com	0/1	ContainerCreating	0	16s

NAME INERS	DESIRED IMAGES	CURRENT SELECTOR	UP-TO-DATE	AVAILABLE	AGE	CONTA
nginx nginx:latest	2 run=nginx	2	2	0	16s	nginx

NAME IMAGES	DESIRED SELECTOR	CURRENT	READY	AGE	CONTAINERS
nginx-6bfb654d7c nginx:latest	2 pod-template-hash=2696210837,run=nginx	2	0	16s	nginx

7.4 A deployment, can be scaled up or down

Copy

```
kubectl scale deploy nginx --replicas=3
```

Output:

```
deployment.extensions "nginx" scaled
```

7.5 Verify the deploy nginx status

Copy

```
kubectl get deploy nginx
```

Output:

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
nginx	3	3	3	3	8m

In a deploy, pods are always controlled by the replica set. However, because the replica set is controlled by the deploy, if we try to scale the replica set instead of the deploy, the deploy will take priority and the number of pods will be reported to the number requested by the deploy.

7.6 Try to scale up the replica set from the previous example to have 5 replicas

Copy

```
replica_set=`kubectl get rs | grep nginx | awk '{print $1}'`  
kubectl scale rs $replica_set --replicas=5
```

Output:

```
replicaset.extensions "nginx-6bfb654d7c" scaled
```

we see the number of pod scaled to 5, according to the request to scale the replica set to 5 pod. After few seconds, the deploy will take priority and remove all new pod created by the scaling the replica set because the desired state, as specified by the deploy is to 3 pods.

7.7 Verify that created pods status

Copy

```
watch kubectl get pods
```


Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx-6bfb654d7c-kglvd	1/1	Running	0	8m
nginx-6bfb654d7c-s8k5f	1/1	Running	0	8m
nginx-6bfb654d7c-x5k72	1/1	Running	0	8m

Note: Wait for couple of minutes to get pods to create and running on **READY** state and then press **<ctrl+c>** to interrupt.

7.8 A deployment also defines the strategy for updates pods

```
strategy:
  rollingUpdate:
    maxSurge: 1
    maxUnavailable: 1
  type: RollingUpdate
```

In the snippet above, we set the update strategy as rolling update.

7.9 Update the pods with a different version of **nginx** image

Copy

```
kubectl set image deploy nginx nginx=nginx:1.13
```

Output:

```
deployment.apps "nginx" image updated
```

7.10 Rollout the deployed nginx status

Copy

```
kubectl rollout status deploy nginx
```

Output:

```
Waiting for rollout to finish: 2 out of 3 new replicas have been updated...

Waiting for rollout to finish: 2 out of 3 new replicas have been updated...

...

...

Waiting for rollout to finish: 1 old replicas are pending termination.
..

...

Waiting for rollout to finish: 2 of 3 updated replicas are available..
.

deployment "nginx" successfully rolled out
```

Note: Wait for couple of minutes to rolled out successfully.

7.11 Check the status of the deploy

Copy

```
kubectl get all -l run=nginx -o wide
```

Output:

NAME	READY	STATUS	RESTARTS	AGE	IP
nginx-59dc74b9-bw5t9 .1.16 pod0-node.onecloud.com	1/1	Running	0	3m	10.244

```
nginx-59dc74b9-sv29f 1/1      Running 0          2m      10.244
.1.17   pod0-node.onecloud.com
```

```
nginx-59dc74b9-xbc96 1/1      Running 0          3m      10.244
.0.10   pod0-node.onecloud.com
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	CONTA
INERS	IMAGES	SELECTOR				
nginx	3	3	3	3	12m	nginx
nginx:1.13	run=nginx					

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS
IMAGES	SELECTOR				
nginx-59dc74b9	3	3	3	3m	nginx
nginx:1.13	pod-template-hash=15873065,run=nginx				
nginx-6bfb654d7c	0	0	0	12m	nginx
nginx:latest	pod-template-hash=2696210837,run=nginx				

Now there is a new replica set now taking control of the pods. This replica set control new pods having image nginx:1.13. The old replica set is still there and can be used in case of downgrade.

7.12 Verify the replicaset status

Copy

```
kubectl get rs
```

Output:

NAME	DESIRED	CURRENT	READY	AGE
nginx-59dc74b9	3	3	3	3m

nginx-6bfb654d7c	0	0	0	13m
------------------	---	---	---	-----

7.13 Verify the pods status

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx-59dc74b9-bw5t9	1/1	Running	0	4m
nginx-59dc74b9-sv29f	1/1	Running	0	3m
nginx-59dc74b9-xbc96	1/1	Running	0	4m

8. Services

A Kubernetes Service is an abstraction which defines a logical set of pods and a policy by which to access them. The set of pods targeted by a Service is usually determined by a label selector. Kubernetes offers a simple Endpoints API that is updated whenever the set of pods in a service changes.

8.1 Create a service for our **nginx webserver**, name as “**nginx-service.yaml**” file

Copy

```
cat > nginx-service.yaml <<EOF

apiVersion: v1

kind: Service

metadata:
  name: nginx

labels:
```

```
    run: nginx
spec:
  selector:
    run: nginx
  ports:
  - protocol: TCP
    port: 8000
    targetPort: 80
  type: ClusterIP
EOF
```

8.2 Create a service for **nginx webserver** service, using **nginx-service.yaml**

Copy

```
kubectl create -f nginx-service.yaml
```

Output:

```
service "nginx" created
```

8.3 Verify the status for nginx service

Copy

```
kubectl get service -l run=nginx
```

Output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
------	------	------------	-------------	---------	-----

nginx	ClusterIP	10.103.148.246	<none>	8000/TCP	7s
-------	-----------	----------------	--------	----------	----

8.4 Describe the nginx service

Copy

```
kubectl describe service nginx
```

Output:

```
Name:          nginx
Namespace:     default
Labels:        run=nginx
Annotations:    <none>
Selector:      run=nginx
Type:          ClusterIP
IP:            10.103.148.246
Port:          <unset> 8000/TCP
TargetPort:    80/TCP
Endpoints:     10.244.0.10:80,10.244.1.16:80,10.244.1.17:80
Session Affinity: None
Events:        <none>
```

The above service is associated to our previous nginx pods. Pay attention to the service selector `run=nginx` field. It tells Kubernetes that all pods with the label `run=nginx` are associated to this service, and should have traffic distributed amongst them. In other words, the service provides an abstraction layer, and it is the input point to reach all of the associated pods.

Pods can be added to the service arbitrarily. Make sure that the label `run=nginx` is associated to any pod we would to bind to the service.

8.5 Create a new pod from the following file without (intentionally) any label

Copy

```
cat > nginx-pod.yaml <<EOF

apiVersion: v1

kind: Pod

metadata:

  name: mynginx

  namespace: default

  labels:

spec:

  containers:

    - name: mynginx

      image: nginx:latest

      ports:

        - containerPort: 80

EOF
```

8.6 Create a new pod, using **nginx-pod.yaml**

Copy

```
kubectl create -f nginx-pod.yaml
```

Output:

```
pod "mynginx" created
```

8.7 Verify the created pods status

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
mynginx	1/1	Running	0	1m
nginx-59dc74b9-bw5t9	1/1	Running	0	6m
nginx-59dc74b9-sv29f	1/1	Running	0	5m
nginx-59dc74b9-xbc96	1/1	Running	0	6m

Note: Wait till pod status changes to **running** state.

8.8 The just created new pod is not still associated to the nginx service, verify the **nginx endpoints**

Copy

```
kubectl get endpoints | grep nginx
```

Output:

nginx	10.244.0.10:80,10.244.1.16:80,10.244.1.17:80	2m
--------------	--	----

8.9 Lets label the new pod with run=nginx label

Copy

```
kubectl label pod mynginx run=nginx
```


Output:

```
pod "mynginx" labeled
```

8.10 We can see a new endpoint is added to the service

Copy

```
kubectl get endpoints | grep nginx
```

Output:

```
nginx          10.244.0.10:80,10.244.1.16:80,10.244.1.17:80 + 1 more...  
3m
```

Any pod in the cluster need for the nginx service will be able to talk with this service by the service address no matter which IP address will be assigned to the nginx pod. Also, in case of multiple nginx pods, the service abstraction acts as load balancer between the nginx pods.

8.11 Create a new pod, name as **busybox.yaml** file

Copy

```
cat > busybox.yaml <<EOF  
  
apiVersion: v1  
  
kind: Pod  
  
metadata:  
  name: busybox  
  namespace: default  
  
spec:  
  containers:  
  - image: busybox
```

```
command:
  - sleep
  - "3600"

imagePullPolicy: IfNotPresent

name: busybox

restartPolicy: Always

EOF
```

8.12 We'll use this pod to address the nginx service

Copy

```
kubectl create -f busybox.yaml

pod "busybox" created
```

8.13 Verify **busybox** pod is in running state:

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
busybox	1/1	Running	0	1m
mynginx	1/1	Running	0	3m
nginx-59dc74b9-bw5t9	1/1	Running	0	9m
nginx-59dc74b9-sv29f	1/1	Running	0	8m

nginx-59dc74b9-xbc96	1/1	Running	0	9m
----------------------	-----	---------	---	----

Note: Wait till pod status changes to **running** state.

8.14 Capture the **service_ip** from the **nginx** service

Copy

```
service_ip=`kubectl get svc | grep nginx | awk '{print $3}'`  
  
echo $service_ip
```

Output:

```
10.103.148.246
```

8.15 Download the busybox application with assigned service_ip

Copy

```
kubectl exec -it busybox -- wget -O - $service_ip:8000  
  
Connecting to 10.254.105.187:8000 (10.254.105.187:8000)  
  
<!DOCTYPE html>  
  
<html>  
  
<head>  
  
<title>Welcome to nginx!</title>  
  
</head>  
  
<body>  
  
<h1>Welcome to nginx!</h1>  
  
<p>If you see this page, the nginx web server is successfully installed and
```

```
working. Further configuration is required.</p>
```

```
...
```

```
</body>
```

```
</html>
```

```
- 100% |*****| 612 0:00:00 ETA
```

8.16 Delete the deployed nginx pod

Copy

```
kubectl delete deploy nginx
```

Output:

```
deployment.extensions "nginx" deleted
```

8.17 Delete the service nginx

Copy

```
kubectl delete svc nginx
```

Output:

```
service "nginx" deleted
```

8.18 Cleanup the other remaining pods by running the below command:

Copy

```
kubectl delete pods --all
```

Output:

```
pod "busybox" deleted
```

```
pod "mynginx" deleted
```

In kubernetes, the service abstraction acts as stable endpoint for any application, no matter which is the IP address of the pod(s) running that application. We can destroy all nginx pods and recreate but service will be always the same IP address and port.

9. Daemons

A Daemon Set is a controller type ensuring each node in the cluster runs a pod. As new node is added to the cluster, a new pod is added to the node. As the node is removed from the cluster, the pod running on it is removed and not scheduled on another node. Deleting a Daemon Set will clean up all the pods it created.

9.1 The configuration file **nginx-daemon-set.yaml** defines a daemon set for the nginx application

Copy

```
cat > nginx-daemon-set.yaml << EOF

apiVersion: extensions/v1beta1

kind: DaemonSet

metadata:

  labels:

    run: nginx

  name: nginx-ds

  namespace:

spec:

  selector:

    matchLabels:
```

```
    run: nginx-ds

template:

  metadata:

    labels:

      run: nginx-ds

  spec:

    containers:

      - image: nginx:latest

        imagePullPolicy: Always

        name: nginx

        ports:

          - containerPort: 80

            protocol: TCP

        dnsPolicy: ClusterFirst

        restartPolicy: Always

        terminationGracePeriodSeconds: 30

EOF
```

9.2 Create a daemon set using **nginx-daemon-set.yaml** file

Copy

```
kubectl create -f nginx-daemon-set.yaml
```

Output:

```
daemonset.extensions "nginx-ds" created
```

9.3 Verify the daemon-set status

Copy

```
kubectl get ds nginx-ds -o wide
```

NAME SELECTOR	DESIRED AGE	CURRENT CONTAINERS	READY IMAGES	UP-TO-DATE	AVAILABLE	NODE SELECTOR
nginx-ds 1m	2 nginx	2 nginx:latest	2 run=nginx-ds	2	2	

Note: Wait till the ready state changes to “2” .

9.4 There are exactly two pods since we have two nodes

Copy

```
kubectl get pods -o wide
```

Output:

NAME NODE	READY	STATUS	RESTARTS	AGE	IP
nginx-ds-kts84 pod38-node.onecloud.com	1/1	Running	0	1m	10.244.1.20
nginx-ds-mtt9k pod38-master.onecloud.com	1/1	Running	0	1m	10.244.0.11

Note: Wait till pod status changes to **running** state.

and each pod is running on a different node.

If you delete a node from the cluster, the running pod on it is removed and not scheduled on other nodes as happens with other types of controllers

Adding back the node to the cluster, you will see a new pod scheduled on that node.

9.5 It is possible to exclude some nodes from the daemon set by forcing the node selector.

```
...

spec:

  containers:

    - image: nginx:latest

      imagePullPolicy: Always

      name: nginx

      ports:

        - containerPort: 80

          protocol: TCP

  ...

  nodeSelector:

    kubernetes.io/hostname: spare.onecloud.com
```

9.6 Cleanup the daemonset nginx

Copy

```
kubectl delete ds nginx-ds
```

Output:

```
daemonset.extensions "nginx-ds" deleted
```