

Lab 7 – Configure Persistent Storage for Kubernetes

Introduction

In this lab, you will learn how to create a persistent volume, persistent volume claim and also reclaim the volume policy and create the storage classes.

Containers are ephemeral, meaning the container file system only lives as long as the container does. Volumes are simplest way to achieve data persistence. In Kubernetes, a more flexible and powerful model is available.

This model is based on the following abstractions:

- **PersistentVolume:** it models shared storage that has been provisioned by the cluster administrator. It is a resource in the cluster just like a node is a cluster resource. Persistent volumes are like standard volumes, but having a lifecycle independent of any individual pod. Also they hide to the users the details of the implementation of the storage, e.g. NFS, iSCSI, or other cloud storage systems.
- **PersistentVolumeClaim:** it is a request for storage by a user. It is similar to a pod. Pods consume node resources and persistent volume claims consume persistent volume objects. As pods can request specific levels of resources like CPU and memory, volume claimers can request the access modes like read-write or read-only and storage capacity.

Kubernetes provides two different ways to provisioning storage:

- **Manual Provisioning:** the cluster administrator has to manually make calls to the storage infrastructure to create persistent volumes and then users need to create volume claims to consume storage volumes.
- **Dynamic Provisioning:** storage volumes are automatically created on-demand when users claim for storage avoiding the cluster administrator to pre-provision storage.

1. Local Persistent Volumes

1.1 Make sure you, logged in as “**root**” user on pod0-spare.origin.com host

Copy

```
ssh root@pod0-master.origin.com
```

1.2 Create a persistent volume as “**local-persistent-volume-recycle.yaml**” file:

Copy

```
cat > local-persistent-volume-recycle.yaml <<EOF

kind: PersistentVolume

apiVersion: v1

metadata:

  name: local-pv

  labels:

    type: local

spec:

  storageClassName: ""

  capacity:

    storage: 2Gi

  accessModes:

    - ReadWriteOnce

  hostPath:

    path: "/data"

  persistentVolumeReclaimPolicy: Recycle

EOF
```

The configuration file specifies that the volume is at **/data** on the the cluster’s node. The volume type is **hostPath** meaning the volume is local to the host node. The configuration also specifies a size of 2GB and the access mode of **ReadWriteOnce**, meanings the volume can be mounted as read write by a single pod at time.

The reclaim policy is **Recycle** meaning the volume can be used many times. It defines the Storage Class name manual for the persistent volume, which will be used to bind a claim to this volume.

1.3 Run the below command to create local persistent volume.

Copy

```
kubectl create -f local-persistent-volume-recycle.yaml
```

Output:

```
persistentvolume "local-pv" created
```

1.4 Verify the persistent volume status

Copy

```
kubectl get pv
```

Output:

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
local-pv	2Gi	RWO	Recycle	Available	
37s					

1.5 Create a persistent volume claim as “**volume-claim.yaml**” file:

Copy

```
cat > volumeclaim-pvc.yaml <<EOF

kind: PersistentVolumeClaim

apiVersion: v1

metadata:

  name: volumeclaim-pvc
```

```
spec:

  storageClassName: ""

  accessModes:

    - ReadWriteOnce

  resources:

    requests:

      storage: 1Gi

EOF
```

Note: The claim is for 1GB of space where the the volume is 2GB. The claim will bound any volume meeting the minimum requirements specified into the claim definition.

1.6 Create the persistent volume claim

Copy

```
kubectl create -f volumeclaim-pvc.yaml
```

Output:

```
persistentvolumeclaim "volumeclaim-pvc" created
```

1.7 Verify that the persistent volume status is in “bound”

Copy

```
kubectl get pv
```

Output:

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			

localpv	2Gi	RWO	Recycle	Bound	default
/volumeclaim-pvc			4m		

1.8 Verify that the persistent volume claim status is in **Bound**

Copy

```
kubect1 get pvc
```

Output:

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAG
ECLASS	AGE				
volumeclaim-pvc	Bound	localpv	2Gi	RWO	
9s					

1.9 Create a **nginx-pod-pvc.yaml** configuration file for a nginx pod using the above claim for its html content directory

Copy

```
cat > nginx-pod-pvc.yaml <<EOF

---

kind: Pod

apiVersion: v1

metadata:

  name: nginx

  namespace: default

  labels:

spec:
```

```
containers:

  - name: nginx

    image: nginx:latest

    ports:

      - containerPort: 80

        name: "http-server"

    volumeMounts:

      - mountPath: "/usr/share/nginx/html"

        name: html

volumes:

  - name: html

    persistentVolumeClaim:

      claimName: volumeclaim-pvc

EOF
```

Note: The pod configuration file specifies a persistent volume claim, but it does not specify a persistent volume. From the pod point of view, the claim is the volume. Please note that a claim must exist in the same namespace as the pod using the claim.

1.10 Create the nginx pod from persistent Volume claim

Copy

```
kubectl create -f nginx-pod-pvc.yaml
```

Output:

```
pod "nginx" created
```

1.11 Accessing the nginx will return **403 Forbidden** since there are no html files to serve in the data volume

Copy

```
kubectl get pod
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	1m

Note: Wait till pod changes to **running** state.

1.12 Verify the nginx pod status and get pod IP's

Copy

```
kubectl get pod nginx -o yaml | grep IP
```

Sample output:

```
hostIP: 10.1.64.246
```

```
podIP: 10.244.1.36
```

1.13 Capture the pod ip's from the nginx server

Copy

```
pod_ip=`kubectl describe pod nginx | grep IP | cut -d " " -f 12`  
  
echo $pod_ip
```

Sample output:

```
10.244.1.36
```

1.14 Try to access the nginx application via `pod_ip` , but there are no html files to serve in the data volume

Copy

```
curl http://$pod_ip:80
```

Output:

```
403 Forbidden
```

1.15 Let's populate the data volume by switching to worker node **pod0-node** .

Copy

```
ssh root@pod0-master.origin.com
```

Copy

```
echo "Welcome to $(hostname)" > /data/index.html
```

1.16 Exit from **pod0-master.origin.com** host:

Copy

```
exit
```

Note:Run the below commands on **pod0-master.origin.com**.

1.17 Now try again to access the nginx application :

Copy

```
curl http://$pod_ip:80
```

Output:


```
Welcome to pod40-node.onecloud.com
```

1.18 To test the persistence of the volume and related claim, delete the pod and recreate it

Copy

```
kubectl delete pod nginx
```

Output:

```
pod "nginx" deleted
```

1.19 Create a nginx pod for persistence volume claim

Copy

```
kubectl create -f nginx-pod-pvc.yaml
```

Output:

```
pod "nginx" created
```

1.20 Verify the nginx pod status is in **“running”**:

Copy

```
kubectl get pod nginx
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	1m

Note: Wait till pod changes to **running** state.

1.21 Locate the IP of the new nginx pod and try to access it

Copy

```
pod_ip=`kubectl describe pod nginx | grep IP | cut -d " " -f 12`  
  
echo $pod_ip
```

Sample output:

```
10.244.1.37
```

1.22 Access the nginx application

Copy

```
curl http://$pod_ip
```

Output:

```
Welcome to pod40-node.onecloud.com
```

2. Volume state

When a pod claims for a volume, the cluster inspects the claim to find the volume meeting claim requirements and mounts that volume for the pod. Once a pod has a claim and that claim is bound, the bound volume belongs to the pod.

A volume will be in one of the following state:

- **Available:** a volume that is not yet bound to a claim
- **Bound:** the volume is bound to a claim
- **Released:** the claim has been deleted, but the volume is not yet available
- **Failed:** the volume has failed

The volume is considered released when the claim is deleted, but it is not yet available for another claim. Once the volume becomes available again then it can bound to another other claim.

2.1 Delete the pod

Copy

```
kubectl delete pod nginx
```

Output:

```
pod "nginx" deleted
```

2.2 Delete the persistent volume claim

Copy

```
kubectl delete pvc volumeclaim-pvc
```

Output:

```
persistentvolumeclaim "volumeclaim-pvc" deleted
```

2.3 Verify the status of the volume

Copy

```
kubectl get pv local-pv
```

Output:

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
localpv	2Gi	RWO	Recycle	Released	default/volumeclaim-pvc
			9m		

3. Volume Reclaim Policy

When deleting a claim, the volume becomes available to other claims only when the volume claim policy is set to Recycle. Volume claim policies currently supported are:

A volume will be in one of the following state:

- **Retain:** the content of the volume still exists when the volume is unbound and the volume is released
- **Recycle:** the content of the volume is deleted when the volume is unbound and the volume is available
- **Delete:** the content and the volume are deleted when the volume is unbound.

Please note that, currently, only NFS and HostPath support recycling.

When the policy is set to Retain the volume is released but it is not yet available for another claim because the previous claimant's data are still on the volume.

3.1 Define a persistent volume using **local-persistent-volume-retain.yaml** configuration file

Copy

```
cat > local-persistent-volume-retain.yaml <<EOF

kind: PersistentVolume

apiVersion: v1

metadata:

  name: local-retain

  labels:

    type: local

spec:

  storageClassName: ""

  capacity:

    storage: 2Gi

  accessModes:

    - ReadWriteOnce

  hostPath:
```

```
    path: "/data"

    persistentVolumeReclaimPolicy: Retain
EOF
```

3.2 Create the persistent volume and the claim

Copy

```
kubectl create -f local-persistent-volume-retain.yaml

kubectl create -f volumeclaim-pvc.yaml
```

Output:

```
persistentvolume "local-retain" created

persistentvolumeclaim "volumeclaim-pvc" created
```

3.3 Create the nginx pod for persistent volume claim

Copy

```
kubectl create -f nginx-pod-pvc.yaml
```

Output:

```
pod "nginx" created
```

3.4 Verify the pod status is in “running”

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	2m

Note: Wait till pod changes to **Running** status, before proceeding to next step .

3.5 Login to the pod using the claim and create some data on the volume

Copy

```
kubectl exec -it nginx bash
```

Output:

```
root@nginx:/#
```

3.6 Add data to index.html file

Copy

```
echo "Hello World" > /usr/share/nginx/html/index.html
```

3.7 Exit from the nginx server

Copy

```
exit
```

3.8 Verify the status of the volume

Copy

```
kubectl get pv
```

Output:

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS
CLAIM		STORAGECLASS	REASON AGE	

local-pv 10m	2Gi	RWO	Recycle	Available
local-retain default/volumeclaim-pvc	2Gi	RWO	Retain	Bound
			3m	

We see the volume remain in the released status and not becomes available since the reclaim policy is set to Retain.

3.9 Now login to the worker node and check data are still there.

An administrator can manually reclaim the volume by deleting the volume and creating a another one.

4. Manual volumes provisioning

In this section we're going to use a Network File System storage backend for manual provisioning of shared volumes. Main limit of local storage for container volumes is that storage area is tied to the host where it resides. If kubernetes moves the pod from another host, the moved pod is no more to access the data since local storage is not shared between multiple hosts of the cluster. To achieve a more useful storage backend we need to leverage on a shared storage technology like NFS.

4.1 Create a NFS share directory on master node as we do not have a dedicated NFS server for this demo.

Copy

```
mkdir /mnt/nfs
```

4.2 Add entry to **/etc/exports** to describe NFS share

Copy

```
cat >> /etc/exports <<EOF

/mnt/nfs *(rw,async,no_root_squash)

EOF
```

4.3 Enable and start NFS server daemon

Copy

```
systemctl enable nfs-server  
  
systemctl start nfs-server
```

Output:

```
Created symlink from /etc/systemd/system/multi-user.target.wants/nfs-s  
erver.service to /usr/lib/systemd/system/nfs-server.service.
```

4.4 Verify the NFS share is available

Copy

```
exportfs
```

Output:

```
/mnt/nfs          <world>
```

4.5 Define a persistent volume as in the **nfs-persistent-volume.yaml** configuration file

Copy

```
cat > nfs-persistent-volume.yaml <<EOF  
  
apiVersion: v1  
  
kind: PersistentVolume  
  
metadata:  
  name: nfs-volume  
  
spec:  
  storageClassName: ""  
  
  capacity:
```



```
    storage: 1Gi

    accessModes:

    - ReadWriteOnce

    nfs:

        path: "/mnt/nfs"

        server: pod0-master

    persistentVolumeReclaimPolicy: Recycle
EOF
```

4.6 Create the nfs persistent volume

Copy

```
kubectl create -f nfs-persistent-volume.yaml
```

Output:

```
persistentvolume "nfs-volume" created
```

4.7 Verify the persistent volume nfs-volume status

Copy

```
kubectl get pv nfs-volume -o wide
```

Output:

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS
CLAIM		STORAGECLASS	REASON AGE	
nfs-volume	1Gi	RWO	Recycle	Available
7s				

4.8 Create a persistent volume claim using **volume-claim.yaml**

Copy

```
cat > volume-claim.yaml <<EOF

kind: PersistentVolumeClaim

apiVersion: v1

metadata:

  name: volume-claim

spec:

  storageClassName: ""

  accessModes:

    - ReadWriteOnce

  resources:

    requests:

      storage: 1Gi

EOF
```

Copy

```
kubectl create -f volume-claim.yaml
```

Output:

```
persistentvolumeclaim "volume-claim" created
```

4.9 Verify volume-claim status is in “bound” status

Copy

```
kubectl get pvc
```

Output:

NAME STORAGECLASS	STATUS AGE	VOLUME	CAPACITY	ACCESS MODES
volume-claim 54m	Bound	nfs-volume	1Gi	RWO
volumeclaim-pvc 59m	Bound	localpv	2Gi	RWO

4.10 Verify persistent volume-claim status

Copy

```
kubectl get pv
```

Output:

NAME CLAIM	CAPACITY	ACCESS MODES STORAGECLASS	RECLAIM POLICY REASON AGE	STATUS
local-pv 13m	2Gi	RWO	Recycle	Available
local-retain default/volumeclaim-pvc	2Gi	RWO	Retain 6m	Bound
nfs-volume default/volume-claim	1Gi	RWO	Recycle 1m	Bound

4.11 Now we are going to create more nginx pods using the same claim.

Create the **nginx-pvc-template.yaml** template for a nginx application having the html content folder placed on the shared storage.

Copy

```
cat > nginx-pvc-template.yaml <<EOF
```

```
apiVersion: extensions/v1beta1
```

```
kind: Deployment
```

```
metadata:
```

```
  generation: 1
```

```
  labels:
```

```
    run: nginx
```

```
  name: nginx-pvc
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchLabels:
```

```
      run: nginx
```

```
  strategy:
```

```
    rollingUpdate:
```

```
      maxSurge: 1
```

```
      maxUnavailable: 1
```

```
    type: RollingUpdate
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
    run: nginx

spec:

  containers:

    - image: nginx:latest

      imagePullPolicy: IfNotPresent

      name: nginx

      ports:

        - containerPort: 80

          protocol: TCP

          name: "http-server"

      volumeMounts:

        - mountPath: "/usr/share/nginx/html"

          name: html

  volumes:

    - name: html

      persistentVolumeClaim:

        claimName: volume-claim

  dnsPolicy: ClusterFirst

  restartPolicy: Always
```

EOF

The template above defines a nginx application based on a nginx deploy of 3 replicas. The nginx application requires a shared volume for its html content. The application does not have to deal with complexity of setup and admin an NFS share.

4.12 Deploy the application

Copy

```
kubectl create -f nginx-pvc-template.yaml
```

Output:

```
deployment.extensions "nginx-pvc" created
```

4.13 Check all pods are up and running

Copy

```
kubectl get pods -o wide
```

Sample output:

NAME IP	NODE	READY	STATUS	RESTARTS	AGE
nginx 10.244.1.5	pod38-node.onecloud.com	1/1	Running	0	7m
nginx-pvc-5846b844f5-2tmnk 10.244.1.6	pod38-node.onecloud.com	1/1	Running	0	1m
nginx-pvc-5846b844f5-cnq4c 10.244.0.3	pod38-node.onecloud.com	1/1	Running	0	1m
nginx-pvc-5846b844f5-dpfsr 10.244.1.7	pod38-node.onecloud.com	1/1	Running	0	1m

Note: Wait till the pod status changes to “**running**” status. **Re-run** the above command to verify of it.

4.14 Login to one of these pods and add html content

Copy

```
kubectl get pods -o wide | grep "nginx-pvc" | awk '{print $1}'
```

Sample output:

```
nginx-pvc-5846b844f5-2tmnk
```

```
nginx-pvc-5846b844f5-cnq4c
```

```
nginx-pvc-5846b844f5-dpfsr
```

4.15 Capture the pod_name from the created nginx persistent volume claim

Copy

```
pod_name=`kubectl get pods -o wide | grep "nginx-pvc" | awk '{print $1}' | head -1`
```

```
echo $pod_name
```

Sample output:

```
nginx-pvc-5846b844f5-2tmnk
```

4.16 Try to access the nginx persistent volume claim server and html content to **index.html**

Copy

```
kubectl exec -it $pod_name bash
```

Output:

```
root@nginx-pvc-5846b844f5-2tmnk:/#
```

Copy

```
echo 'Hello from NFS!' > /usr/share/nginx/html/index.html
```

Copy

```
exit
```

4.17 Since all three pods mount the same shared folder on the NFS, the just created html content is placed on the NFS share and it is accessible from any of the three pods

Copy

```
kubectl get pod -o wide | grep nginx-pvc | awk '{print $6}'
```

Sample output:

```
10.32.0.3
```

```
10.44.0.1
```

```
10.44.0.2
```

4.18 Capture the pod ip from the nginx pvc

Copy

```
pod_ip=`kubectl get pod -o wide | grep nginx-pvc | awk '{print $6}' |  
head -1`
```

```
echo $pod_ip
```

Sample output:

```
10.32.0.3
```

Copy

```
curl http://$pod_ip
```


Output:

```
Hello from NFS!
```

5. Cleanup

5.1 Delete all the pods

Copy

```
kubectl delete pods --all
```

Sample output:

```
pod "nginx" deleted  
  
pod "nginx-pvc-5846b844f5-2tmnk" deleted  
  
pod "nginx-pvc-5846b844f5-cnq4c" deleted  
  
pod "nginx-pvc-5846b844f5-dpfsr" deleted
```