

Lab 8 – Multiple-tier applications

Containers

Introduction

In this lab, you will learn how to build and deploy a simple, multi-tier web application using Kubernetes and Docker.

Create the Redis master pod

The guestbook application uses Redis to store its data. It writes its data to a Redis master instance and reads data from multiple Redis slave instances

Creating the Redis Master Deployment

The manifest file, included below, specifies a Deployment controller that runs a single replica Redis master Pod.

- Launch a terminal window in the directory you downloaded the manifest files.
- Apply the Redis Master Deployment from the yaml file

1. Deploying an Multi-tier web application

1.1 Create the Redis Master Deployment from the **redis-master-deployment.yaml** file:

Copy

```
cat > redis-master-deployment.yaml <<EOF

apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2

kind: Deployment

metadata:

  name: redis-master

spec:

  selector:
```

```
matchLabels:

  app: redis

  role: master

  tier: backend

replicas: 1

template:

  metadata:

    labels:

      app: redis

      role: master

      tier: backend

  spec:

    containers:

      - name: master

        image: k8s.gcr.io/redis:e2e # or just image: redis

        resources:

          requests:

            cpu: 100m

            memory: 100Mi

        ports:

          - containerPort: 6379
```

EOF

1.2 Create the **redis-master-deployment.yaml** file

Copy

```
kubectl apply -f redis-master-deployment.yaml
```

1.3 Query the list of pods to verify that the redis master pod is running:

Copy

```
kubectl get pods
```

Sample output:

NAME	READY	STATUS	RESTARTS	AGE
redis-master-1068406935-3lswp	1/1	Running	0	28s

Note: Wait till pod changes to **Running** status, before proceeding to next step .

1.4 Run the following command to view the logs from the Redis Master Pod:

Capture the redis_pod name from the pod status

Copy

```
redis_pod=`kubectl get pods | grep redis-master | awk '{print $1}'`  
echo $redis_pod
```

Sample output:

```
redis-master-55db5f7567-qmq4v
```

1.5 Verify the logs from assigned variable redis_pod

Copy

```
kubectl logs -f $redis_pod
```

Output:

[illegible]

```
[1] 22 Apr 11:45:17.967 # Server started, Redis version 2.8.19
```

```
[1] 22 Apr 11:45:17.967 # WARNING you have Transparent Huge Pages (THP) support enabled in your kernel. This will create latency and memory usage issues with Redis. To fix this issue run the command 'echo never > /sys/kernel/mm/transparent_hugepage/enabled' as root, and add it to your /etc/rc.local in order to retain the setting after a reboot. Redis must be restarted after THP is disabled.
```

```
[1] 22 Apr 11:45:17.967 # WARNING: The TCP backlog setting of 511 cannot be enforced because /proc/sys/net/core/somaxconn is set to the lower value of 128.
```

```
[1] 22 Apr 11:45:17.967 * The server is now ready to accept connections on port 6379
```

Note: Press <ctrl+c> to exit.

2. Creating the Redis Master Service

2.1 The guestbook application needs to communicate to the Redis master to write its data. You need to apply a Service to proxy the traffic to the Redis master Pod. A Service defines a policy to access the Pods.

Copy

```
cat > redis-master-service.yaml <<EOF

apiVersion: v1

kind: Service

metadata:

  name: redis-master

  labels:

    app: redis

    role: master

    tier: backend
```

```
spec:

  ports:
    - port: 6379

      targetPort: 6379

  selector:

    app: redis

    role: master

    tier: backend

EOF
```

Note: This manifest file creates a Service named redis-master with a set of labels that match the labels previously defined, so the Service routes network traffic to the Redis master Pod.

2.2 Create the **redis-master-service.yaml** file

Copy

```
kubectl apply -f redis-master-service.yaml
```

Output:

```
service "redis-master" created
```

2.3 Query the list of Services to verify that the Redis Master Service is running:

Copy

```
kubectl get service
```

Output:

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
------	------------	-------------	---------	-----

kubernetes 2h	ClusterIP	10.96.0.1	<none>	443/TCP
redis-master 4s	ClusterIP	10.108.95.60	<none>	6379/TCP

Start up the Redis Slaves

Although the Redis master is a single pod, you can make it highly available to meet traffic demands by adding replica Redis slaves.

3. Creating the Redis Slave Deployment

Deployments scale based off of the configurations set in the manifest file. In this case, the Deployment object specifies two replicas.

If there are not any replicas running, this Deployment would start the two replicas on your container cluster. Conversely, if there are more than two replicas are running, it would scale down until two replicas are running.

3.1 Create the slave deploy, from **redis-slave-deployment.yaml** file

Copy

```
cat > redis-slave-deployment.yaml <<EOF

apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2

kind: Deployment

metadata:

  name: redis-slave

spec:

  selector:

    matchLabels:

      app: redis
```

```
    role: slave

    tier: backend

replicas: 2

template:

  metadata:

    labels:

      app: redis

      role: slave

      tier: backend

  spec:

    containers:

      - name: slave

        image: gcr.io/google_samples/gb-redisslave:v1

        resources:

          requests:

            cpu: 100m

            memory: 100Mi

        env:

          - name: GET_HOSTS_FROM

            value: dns

            # Using `GET_HOSTS_FROM=dns` requires your cluster to
```



```

        # provide a dns service. As of Kubernetes 1.3, DNS is a built-in
        # service launched automatically. However, if the cluster you are using
        # does not have a built-in DNS service, you can instead
        # access an environment variable to find the master
        # service's host. To do so, comment out the 'value: dns' line above, and
        # uncomment the line below:
        # value: env

ports:
  - containerPort: 6379

EOF

```

3.2 Apply the Redis Slave Deployment from the **redis-slave-deployment.yaml** file:

Copy

```
kubectl apply -f redis-slave-deployment.yaml
```

Output:

```
deployment.apps "redis-slave" created
```

3.3 Query the list of Pods to verify that the Redis Slave Pods are running:

Copy

```
watch kubectl get pods
```

Sample output:

NAME AGE	READY	STATUS	RESTARTS
redis-master-1068406935-3lswp 1m	1/1	Running	0
redis-slave-2005841000-fpvqc 6s	1/1	Running	0
redis-slave-2005841000-phfv9 6s	1/1	Running	0

Note: Wait for couple of minutes to get pod status as **ready** and then press **<ctrl+c>** to interrupt.

4. Creating the Redis Slave Service

The guestbook application needs to communicate to Redis slaves to read data. To make the Redis slaves discoverable, you need to set up a Service. A Service provides transparent load balancing to a set of Pods.

4.1 Create the slave service deployment, from the **redis-slave-service.yaml** file

Copy

```
cat > redis-slave-service.yaml <<EOF

apiVersion: v1

kind: Service

metadata:
  name: redis-slave

  labels:
    app: redis
    role: slave
    tier: backend
```

```
spec:
  ports:
    - port: 6379
  selector:
    app: redis
    role: slave
    tier: backend
EOF
```

4.2 Apply the Redis Slave Service from the following redis-slave-service.yaml file:

Copy

```
kubectl apply -f redis-slave-service.yaml
```

Output:

```
service "redis-slave" created
```

4.3 Query the list of Services to verify that the Redis Slave Service is running:

Copy

```
kubectl get services
```

Output:

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP
2h				

redis-master 5m	ClusterIP	10.108.95.60	<none>	6379/TCP
redis-slave 1m	ClusterIP	10.97.140.193	<none>	6379/TCP

5. Set up and Expose the Guestbook Frontend

The guestbook application has a web frontend serving the HTTP requests written in PHP. It is configured to connect to the **redis-master** Service for write requests and the **redis-slave** service for Read requests.

5.1 Creating the Guestbook Frontend Deployment

Copy

```
cat > frontend-deployment.yaml <<EOF

apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2

kind: Deployment

metadata:
  name: frontend

spec:
  selector:
    matchLabels:
      app: guestbook
      tier: frontend

  replicas: 3

  template:
    metadata:
```

```

labels:

  app: guestbook

  tier: frontend

spec:

  containers:

  - name: php-redis

    image: gcr.io/google-samples/gb-frontend:v4

    resources:

      requests:

        cpu: 100m

        memory: 100Mi

    env:

    - name: GET_HOSTS_FROM

      value: dns

      # Using `GET_HOSTS_FROM=dns` requires your cluster to
      # provide a dns service. As of Kubernetes 1.3, DNS is a built-in
      # service launched automatically. However, if the cluster you are using
      # does not have a built-in DNS service, you can instead
      # access an environment variable to find the master
      # service's host. To do so, comment out the 'value: dns' line above, and

```

```
# uncomment the line below:

# value: env

ports:

- containerPort: 80

EOF
```

5.2 Apply the frontend Deployment from the following frontend-deployment.yaml file:

Copy

```
kubectl apply -f frontend-deployment.yaml
```

Output:

```
deployment.apps "frontend" created
```

5.3 Query the list of Pods to verify that the three frontend replicas are running:

Copy

```
watch kubectl get pods -l app=guestbook -l tier=frontend
```

Sample output:

NAME	READY	STATUS	RESTARTS	AGE
frontend-3823415956-dsvc5	1/1	Running	0	54s
frontend-3823415956-k22zn	1/1	Running	0	54s
frontend-3823415956-w9gbt	1/1	Running	0	54s

Note: Wait for couple of minutes to get pods to create and running on **READY** state and then press **<ctrl+c>** to interrupt.

a. Creating the Frontend Service

The **redis-slave** and **redis-master** Services you applied are only accessible within the container cluster because the default type for a Service is ClusterIP. **ClusterIP** provides a single IP address for the set of Pods the Service is pointing to. This IP address is accessible only within the cluster.

If you want guests to be able to access your guestbook, you must configure the frontend Service to be externally visible, so a client can request the Service from outside the container cluster. Minikube can only expose Services through **NodePort**.

5.3 Create frontend-service.yaml

Copy

```
cat > frontend-service.yaml <<EOF

apiVersion: v1

kind: Service

metadata:

  name: frontend

  labels:

    app: guestbook

    tier: frontend

spec:

  # comment or delete the following line if you want to use a LoadBalancer

  type: NodePort

  # if your cluster supports it, uncomment the following to automatically
  # create an external load-balanced IP for the frontend service.

  # type: LoadBalancer

  ports:
```

```
- port: 80

selector:

  app: guestbook

  tier: frontend

EOF
```

5.4 Apply the frontend Service from the following frontend-service.yaml file

Copy

```
kubectl apply -f frontend-service.yaml
```

Output:

```
service "frontend" created
```

5.5 Query the list of Services to verify that the frontend Service is running:

Copy

```
kubectl get services
```

Output

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
frontend 4s	NodePort	10.97.74.175	<none>	80:31288/TCP
kubernetes 2h	ClusterIP	10.96.0.1	<none>	443/TCP
redis-master 10m	ClusterIP	10.108.95.60	<none>	6379/TCP

redis-slave 6m	ClusterIP	10.97.140.193	<none>	6379/TCP
-------------------	-----------	---------------	--------	----------

b. Viewing the Frontend Service via NodePort

5.6 Capture and take a note of node_port from the frontend service

Copy

```
node_port=`kubectl get services | grep frontend | awk '{print $5}' | cut -d ":" -f 2 | cut -d "/" -f 1`  
  
echo $node_port
```

Sample output:

```
30195
```

5.7 Copy the below URL, and load the page in your browser to view your guestbook.

Copy

```
http://localhost:node_port
```

Note: Make sure the port-forwarding is configured, Ex: **30195**.

6. Scale the Web Frontend

Scaling up or down is easy because your servers are defined as a Service that uses a Deployment controller.

6.1 Run the following command to scale up the number of frontend Pods:

Copy

```
kubectl scale deployment frontend --replicas=5
```

Output:

```
deployment.extensions "frontend" scaled
```

6.2 Query the list of Pods to verify the number of frontend Pods running:

Copy

```
kubectl get pods
```

Sample output:

NAME	READY	STATUS	RESTARTS	AGE
frontend-5c548f4769-4lv9g	1/1	Running	0	37s
frontend-5c548f4769-7zdcB	1/1	Running	0	6m
frontend-5c548f4769-cqngw	1/1	Running	0	6m
frontend-5c548f4769-wb2st	1/1	Running	0	6m
frontend-5c548f4769-z8r2n	1/1	Running	0	37s
nginx	1/1	Running	0	35m
redis-master-55db5f7567-qmq4v	1/1	Running	0	16m
redis-slave-584c66c5b5-fm5cd	1/1	Running	0	11m
redis-slave-584c66c5b5-jj7p9	1/1	Running	0	11m

6.3 Run the following command to scale down the number of frontend Pods:

Copy

```
kubectl scale deployment frontend --replicas=2
```

Output:

```
deployment.extensions "frontend" scaled
```

6.4 Query the list of Pods to verify the number of frontend Pods running:

Copy

```
kubectl get pods
```

Sample output:

NAME	READY	STATUS	RESTARTS	AGE
frontend-5c548f4769-cqngw	1/1	Running	0	7m
frontend-5c548f4769-wb2st	1/1	Running	0	7m
nginx	1/1	Running	0	37m
redis-master-55db5f7567-qmq4v	1/1	Running	0	17m
redis-slave-584c66c5b5-fm5cd	1/1	Running	0	12m
redis-slave-584c66c5b5-jj7p9	1/1	Running	0	12m

7. Cleanup

Deleting the Deployments and Services also deletes any running Pods. Use labels to delete multiple resources with one command.

7.1 Run the following commands to delete all Pods, Deployments, and Services.

Copy

```
kubectl delete service -l app=redis  
  
kubectl delete service -l app=guestbook  
  
kubectl delete deploy frontend redis-master redis-slave
```

Output:

```
deployment "redis-master" deleted  
  
deployment "redis-slave" deleted  
  
service "redis-master" deleted  
  
service "redis-slave" deleted  
  
deployment "frontend" deleted
```

```
service "frontend" deleted
```

7.2 Query the list of Pods to verify that no Pods are running:

Copy

```
kubectl get pods
```

Output:

```
No resources found.
```