

# Lab 11 – Configure Quotas and Limits

## Quotas and Limits

Namespaces let different users or teams to share a cluster with a fixed number of nodes. It can be a concern that one team could use more than its fair share of resources. Resource quotas are the tool to address this concern.

A resource quota provides constraints that limit aggregate resource consumption per namespace. It can limit the quantity of objects that can be created in a namespace by type, as well as the total amount of compute resources that may be consumed in that project.

Users create resources in the namespace, and the quota system tracks usage to ensure it does not exceed hard resource limits defined in the resource quota. If creating or updating a resource violates a quota constraint, the request will fail. When quota is enabled in a namespace for compute resources like cpu and memory, users must specify resources consumption, otherwise the quota system rejects pod creation.

### Applying Resource Quotas and Limits

This lab demonstrates a typical setup to control for resource usage in a namespace.

The cluster-admin is operating a cluster on behalf of a user population and the cluster-admin wants to control the amount of resources that can be consumed in a particular namespace to promote fair sharing of the cluster and control cost.

The cluster-admin has the following goals:

- Limit the amount of compute resource for running pods
- Limit the number of persistent volume claims to control access to storage
- Limit the number of load balancers to control cost
- Prevent the use of node ports to preserve scarce resources
- Provide default compute resource requests to enable better scheduling decisions

### Create a namespace

Let's create a new namespace called quota-example:

Copy

```
kubectl config set-context quota-example
```

**Output:**

```
Context "quota-example" created.
```

Copy

```
kubectl config use-context quota-example
```

**Output:**

```
Switched to context "quota-example"
```

Copy

```
kubectl config get-contexts
```

**Output:**

CURRENT NAMESPACE	NAME	CLUSTER	AUTHINFO
	kubernetes-admin@kubernetes	kubernetes	kubernetes-admin
	multi-app		
	-help		

Copy

```
cat > quota.yaml <<EOF

apiVersion: v1

kind: ResourceQuota
```

```
metadata:
```

```
  name: quota
```

```
spec:
```

```
  hard:
```

```
    cpu: "20"
```

```
    memory: 1Gi
```

```
    persistentvolumeclaims: "10"
```

```
    pods: "10"
```

```
    replicationcontrollers: "20"
```

```
    resourcequotas: "1"
```

```
    secrets: "10"
```

```
    services: "5"
```

```
EOF
```

Copy

```
kubectl create -f quota.yaml
```

**Output:**

```
resourcequota "quota" created
```

Copy

```
kubectl describe quota quota
```

**Output:**

Name: project-quota

Namespace: default

Resource	Used	Hard
----------	------	------

-----	----	----
-------	------	------

limits.cpu	0	1
------------	---	---

limits.memory	0	1Gi
---------------	---	-----

Pods	0	10
------	---	----

Name: quota

Namespace: default

Resource	Used	Hard
----------	------	------

-----	----	----
-------	------	------

CPU	0	20
-----	---	----

memory	0	1Gi
--------	---	-----

persistentvolumeclaims	1	10
------------------------	---	----

Pods	0	10
------	---	----

replicationcontrollers	0	20
------------------------	---	----

resourcequotas	2	1
----------------	---	---

secrets	4	10
---------	---	----

services	1	5
----------	---	---

## Applying default resource requests and limits

Pod authors rarely specify resource requests and limits for their pods.

Since we applied a quota to our project, let's see what happens when an end-user creates a pod that has unbounded cpu and memory by creating an nginx container.

Copy

```
kubectl run nginx1 --image=nginx --replicas=1
```

### Output:

```
deployment.apps "nginx1" created
```

Now let's look at the pods that were created.

Copy

```
kubectl get pods
```

### Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx1-65899c769f-kbgj7	1/1	Running	0	59s

Copy

```
kubectl describe deployment nginx1
```

### Output:

```
Name:          nginx
Namespace:     quota-example
CreationTimestamp:  Fri, 13 Apr 2018 14:44:30 +0000
Labels:       run=nginx
```

Annotations: deployment.kubernetes.io/revision=1

Selector: run=nginx

Replicas: 1 desired | 1 updated | 1 total | 1 available  
| 0 unavailable

StrategyType: RollingUpdate

MinReadySeconds: 0

RollingUpdateStrategy: 1 max unavailable, 1 max surge

Pod Template:

Labels: run=nginx

Containers:

nginx:

Image: nginx

Port:

Host Port:

Environment:

Mounts:

Volumes:

Conditions:

Type	Status	Reason
----	-----	-----
Available	True	MinimumReplicasAvailable
Progressing	True	NewReplicaSetAvailable

OldReplicaSets:

NewReplicaSet: nginx-65899c769f (1/1 replicas created)

Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	ScalingReplicaSet	1m	deployment-controller	Scaled up replica set nginx-65899c769f to 1

Copy

```
kubectl describe rs nginx1
```

### Output:

```
Name:          nginx-65899c769f
Namespace:     quota-example
Selector:      pod-template-hash=2145573259,run=nginx
Labels:        pod-template-hash=2145573259
               run=nginx
Annotations:   deployment.kubernetes.io/desired-replicas=1
               deployment.kubernetes.io/max-replicas=2
               deployment.kubernetes.io/revision=1
Controlled By: Deployment/nginx
Replicas:      1 current / 1 desired
Pods Status:   1 Running / 0 Waiting / 0 Succeeded / 0 Failed
```

### Pod Template:

Labels: pod-template-hash=2145573259

run=nginx

### Containers:

nginx:

Image: nginx

Port:

Host Port:

Environment:

Mounts:

Volumes:

### Events:

Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	SuccessfulCreate	1m	replicaset-controller	Created pod: nginx-65899c769f-kbgj7

### Copy

```
cat > limits.yaml <<EOF
```

```
apiVersion: v1
```

```
kind: LimitRange
```

```
metadata:
```

```
  name: limits
```



```
spec:
  limits:
  - default:
      cpu: 200m
      memory: 512Mi
    defaultRequest:
      cpu: 100m
      memory: 256Mi
    type: Container
EOF
```

Copy

```
kubectl create -f limits.yaml
```

**Output:**

```
limitrange "limits" created
```

Copy

```
kubectl describe limits limits
```

**Output:**

```
Name:      limits
Namespace:  quota-example
```

Type	Resource	Min	Max	Default	Request	Default Limit	Max Li
mit/Request	Ratio						

Container	cpu	-	-	100m		200m	-
Container	memory	-	-	256Mi		512Mi	-

Copy

```
kubectl run nginx2 --image=nginx --replicas=1 --requests=cpu=100m,memory=256Mi --limits=cpu=200m,memory=512Mi
```

Copy

```
kubectl get pods
```

Output:

NAME	READY	STATUS	RESTARTS	AGE
nginx1-65899c769f-8hgf7	1/1	Running	0	9m
nginx2-7f4cff6589-gr95w	1/1	Running	0	5m

Copy

```
kubectl describe pod nginx2
```